

High Performance Computing (HPC)

Assignment 06 – Parallel Interpolation

Instructor: Dr. Bhaskar Chaudhury

Teaching Assistant: Libin Varghese, Ayushi Sharma, Kaushik Prajapati

April 1, 2026

General Instructions

- This assignment focuses on **parallelizing** the given interpolation algorithm.
 - You must begin with a **correct serial baseline** and then develop a parallel implementation.
 - You are allowed to use **OpenMP** for parallelization. MPI is not allowed for this assignment.
 - The evaluation will be based on **correctness, scalability, performance, and analysis**.
 - You are encouraged to experiment with **data structures, memory layout, and parallelization strategies**.
 - Submission will be done through a separate Google Form (link will be shared).
-

1 Recap from the Previous Assignment

In Assignment 03, we first introduced the scatter point-to-mesh interpolation problem. In this problem, a set of particles was randomly initialized inside a 1×1 domain, and their values were interpolated onto a structured $N_x \times N_y$ grid. The goal was to study how the execution time changes with different grid sizes and numbers of particles. In this assignment, you are only required to parallelise the interpolation process.

2 Problem Description

Given a set of scattered points in a 2D domain with known function values:

$$S = \{(x_i, y_i, f_i) \mid i = 1, 2, \dots, N\}$$

where (x_i, y_i) are spatial coordinates and f_i are known function values. For this assignment, assume $f_i = 1$. The task is to interpolate these values onto a structured mesh grid of size $M \times M$.

The structured mesh consists of regularly spaced grid points:




Figure 1: Figure of the domain with scattered points

$$G = \{(X_i, Y_j) \mid X_i = i\Delta x, \quad Y_j = j\Delta y, \quad i, j = 0, 1, \dots, M-1\}$$

where Δx and Δy define the uniform grid spacing and $0 \leq x_i \leq X_{max}$ and $0 \leq y_i \leq Y_{max}$. See Fig.1 for reference.

$$\Delta x = \frac{X_{max}}{M}, \quad \Delta y = \frac{Y_{max}}{M}$$

All scattered points lie within a domain of size 1×1 .
Typical use of this scheme is mentioned here.

1. Computer Graphics and 3D Modeling:

- Used to generate smooth surfaces from scattered 3D points, such as those obtained from 3D scanning (e.g., LIDAR or photogrammetry). For example, turning a point cloud of a scanned object into a mesh for rendering or animation.

2. Scientific Visualization:

- In fields like meteorology or oceanography, scattered data points (e.g., temperature or pressure readings from sensors) are interpolated onto a mesh to visualize continuous fields over a region.

3. Finite Element Analysis (FEA):

- Engineers use this in simulations (e.g., structural analysis, fluid dynamics) to map scattered data (like material properties or experimental measurements) onto a structured mesh for numerical solving.




Figure 2: Linear interpolation illustration. Here h_x, h_y correspond to x_i, y_i

4. Medical Imaging:

- Reconstructing surfaces or volumes from scattered data points, such as creating a 3D model of an organ from MRI or CT scan points.

5. Machine Learning and Data Science:

- When dealing with sparse or unevenly sampled data, this scheme can help create a continuous representation for tasks like regression, visualization, or simulation.

3 Linear Interpolation

Linear interpolation is a simple **bilinear interpolation** method in which each scattered point distributes its value to the **four nearest grid points**, proportionally to its relative position inside the grid cell. The procedure consists of:

1. Identify the four nearest grid points (X_i, Y_j) surrounding a scattered point (x_i, y_i) .
2. Compute the weights based on the relative distances of (x_i, y_i) to these grid points.
3. Distribute the value f_i proportionally to these grid points.

For a scattered point (x_i, y_i) located in a grid cell with corners at: (X_i, Y_j) , (X_{i+1}, Y_j) , (X_i, Y_{j+1}) , (X_{i+1}, Y_{j+1}) . The weight factors are computed as:

$$w_{i,j} = (\Delta x - l_x) \times (\Delta y - l_y)$$

$$w_{i+1,j} = l_y \times (\Delta x - l_x)$$

$$w_{i,j+1} = l_x \times (\Delta y - l_y)$$

$$w_{i+1,j+1} = l_x \times l_y$$

where:

$$l_x = x_i - X_i, \quad l_y = y_j - Y_j$$

Since the values of the grid point are not necessary to be stored explicitly, the indices can be calculated on the fly:

$$X_i = \left\lfloor \frac{x_i}{\Delta x} \right\rfloor \times \Delta x, \quad Y_j = \left\lfloor \frac{y_i}{\Delta y} \right\rfloor \times \Delta y$$

The interpolated function values at the structured grid points are accumulated as:

$$\begin{aligned} F(X_i, Y_j) &+ = w_{i,j} \times f_i \\ F(X_{i+1}, Y_j) &+ = w_{i+1,j} \times f_i \\ F(X_i, Y_{j+1}) &+ = w_{i,j+1} \times f_i \\ F(X_{i+1}, Y_{j+1}) &+ = w_{i+1,j+1} \times f_i \end{aligned}$$

For example, the value to be stored on grid point (i, j) depends on the value of weight $w_{i,j}$, which is calculated using the area of the blue region as shown in Fig. 2.

In a serial implementation, interpolation is straightforward and requires two data structures: one for storing scattered points and another for storing interpolated values on a structured mesh. However, in a parallel implementation, special care must be taken to eliminate dependencies. This is because multiple scattered points may fall within the same or adjacent cells, leading to shared grid points and potential race conditions. To improve performance, these dependencies must be carefully managed. **This can be achieved by modifying the data structures or optimizing the parallel implementation to ensure efficient computation. Your final scores will be evaluated based on the leaderboard, and your results will be shared in Google Classroom. Your end goal is to increase the performance of the existing code for a large number of threads.**

4 Implementation Steps

1. **Read input data:** Load scattered points (x_i, y_i, f_i) from *input.bin*. The data in this file is in binary format (to save the space). The domain size is 1 unit $\times$ 1 unit. All particles and grid points should lie within this domain.
2. **Initialize structured mesh:** Define the grid G and set all function values $F(X_i, Y_j)$ to zero.
3. **Perform interpolation:**
 - Loop through each scattered point.
 - Identify the four nearest structured grid points.
 - Compute the weights w .
 - Distribute the function value accordingly.
 - Another way to do this is to Loop through each grid point.
 - Identify all the scattered points associated with the grid point.

- Compute the weights w .
 - Distribute the function value accordingly.
 - This step can be implemented in multiple ways. Think carefully about memory access patterns.
 - First, write a correct serial code and validate results using provided test data.
 - You are free to use any data structure that improves performance.
 - Ensure input and output formats are not altered.
4. Step 3 will be performed 10 times, each time loading a new set of scattered points.
5. **Output the interpolated grid:** Write the structured mesh values to an output file.
6. To check accuracy, run the code using `test_input.bin` and compare with `test_output.txt`.
Performance improvement is important but not at the cost of correctness.
7. Use *input_fileMaker.c* to generate input data with the following configurations:
- (a) $N_x=250$, $N_y=100$, points=0.9 million, Maxiter=10
 - (b) $N_x=250$, $N_y=100$, points=5 million, Maxiter=10
 - (c) $N_x=500$, $N_y=200$, points=3.6 million, Maxiter=10
 - (d) $N_x=500$, $N_y=200$, points=20 million, Maxiter=10
 - (e) $N_x=1000$, $N_y=400$, points=14 million, Maxiter=10

5 Parallelization Objective

The goal of this assignment is to transform the given serial interpolation code into an efficient **parallel implementation**.

Key challenges include:

- Handling **race conditions** due to multiple particles updating the same grid points.
- Designing **efficient memory access patterns**.
- Achieving good **load balancing** across threads/processes.
- Minimizing synchronization overhead.

You are expected to explore different strategies such as:

- Domain decomposition
- Particle decomposition
- Use of atomic operations vs privatization
- Reduction-based approaches

6 Analysis and Questions

1. Write the pseudocode of your parallel implementation in the report along with an illustrative diagram of your implementation.
2. Explain your parallel implementation approach and how you handled race conditions in your code.
3. Plot graphs for each of the five input configurations for speedup vs cores and execution time vs cores (increment cores starting from 2 to 16 by a factor of 2). Ensure no one else is using the machine while running the code to get good performance.
4. Compute the parallel efficiency and comment on the scalability of your implementation.
5. Compare the performance of your parallel code with the serial version and report the maximum speedup achieved.
6. Try to explain the speedup achieved and all the observations using your current HPC knowledge.
7. If given unlimited HPC resources, how would you further optimize the implementation?
8. Discuss any load imbalance or performance bottlenecks observed in your implementation.
9. Suggest at least one alternative approach (data structure or parallelization method) to improve performance.
10. You are required to come up with different data structures or OpenMP-based approaches to increase the performance of the code. Anyone coming up with good results may get bonus marks for this assignment. Leaderboard will be based on performance and efficiency.

Note: Ensure that the implementation correctly interpolates values. The accuracy of your code will be tested against a hidden dataset. Incorrect implementations will not be evaluated for performance.

7 Submission Guidelines

- Submit a PDF report including:
 - Explanation of your implementation approach
 - Performance analysis
 - Answers to analysis questions
- Submit source code in compressed format.
- Submission will be done through a separate Google Form (link will be shared).

8 Code Repository and Remarks

The base implementation and supporting files are available in the shared GitHub repository: [Click here](#)

With so much information and resources at your disposal, sincere effort is expected. **All the best.**